

SpiNNaker2 MAC Array (16x2)

Performance Report

Note: this performance report is based on VGG-19 model with the best mapping strategy.

- **Mapping strategy:**

Convolution layer: operator fusion + data reuse

Fully-connected layer: operator fusion

→ Operator fusion will decrease the size of output activation

→ Data reuse will reduce bandwidth problem of dram

(Next slide has details information for the mapping strategy)

- **Splitting scheme difference between MAC array (16*4) and MAC array (16*2)**

No different

- **Simulator difference between MAC array (16*4) and MAC array (16*2)**

MAC array size is different $16*4 \rightarrow 16*2$;

The prefetching of filter for CONV is different.

For 16*4_MAC_array, each 4 clocks send a fetching request.

For 16*2_MAC_array, each 8 clocks send a fetching request.

Mapping strategy

Convolution layer: operator fusion + data reuse

Operator Fusion

Input-activation into SRAM

Decrease size of input activation

padding

convolution

activation

Decrease size of output activation

pooling

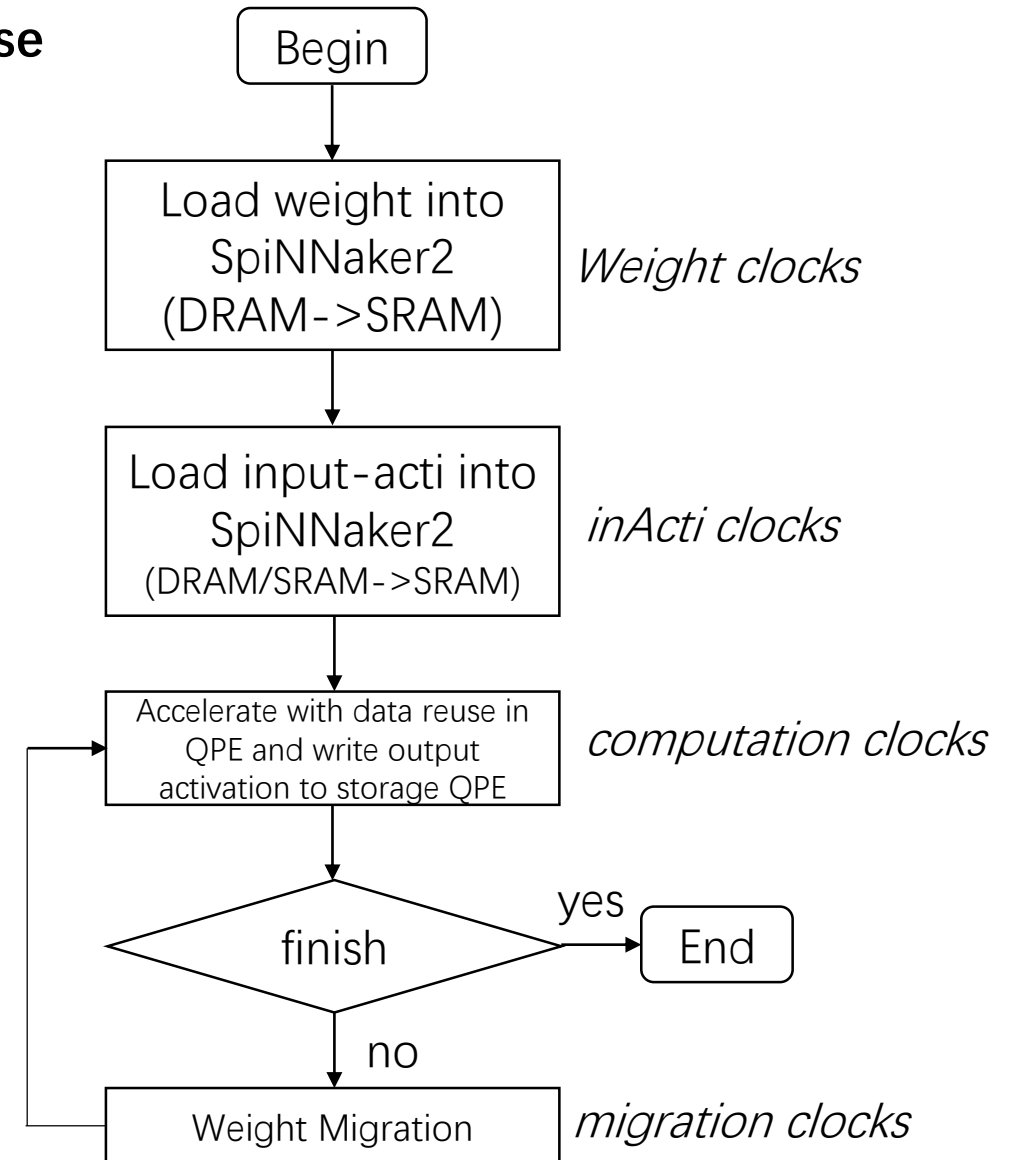
Decrease size of output activation

quantization

output-activation to SRAM

Note that padding, activation, pooling and quantization is not simulated. It is assume to be done immediately (0 clock).

Data reuse



Mapping strategy

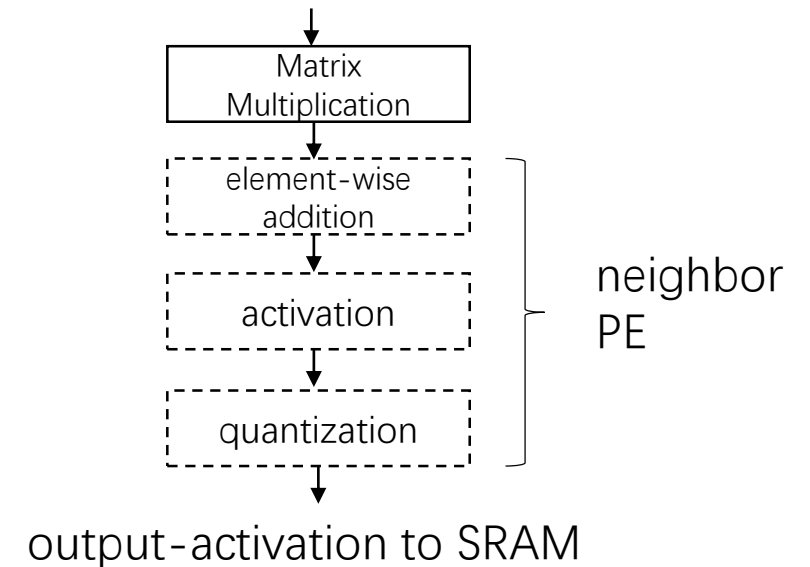
Fully-connected layer: operator fusion

For fully-connected layer, operator fusion and data reuse cannot coexist. As for the convolution, its result (the output activation of CONV will be the input activation of FC) will be stored in QPE. Therefore, data reuse in FC has no meaning, because only input-activation could be reused and due to NoC bandwidth is larger than DRAM bandwidth, reuse input-activation will not increase the system performance. **The whole system performance of FC will be limited by DRAM bandwidth.**

Due to the DRAM bandwidth, most the Pes in SpiNNaker2 will be idle. Theoretically, 8 PEs can fully utilize 4 DRAM.

However, because FC is splitted, some computation (element-wise addition) muss be done ARM core. In order to eliminate this part, some idle PEs (not the PE used for FC) will be used to accelerate the element-wise addition. Moreover, these PEs will also be used to accelerate non-linearity and quantization.

Input-activation and weight into SRAM



Result for VGG-19

MAC array Row = 4, Column = 16

"VGG-CONV_1" : 37982

weight_clock: 497

inActi_clock: 8374

weight_migration: 0

computation: 29111 (including write outActi from SRAM to SRAM)

"VGG-CONV_6" : 214575

weight_clock: 9611

inActi_clock: 23344

weight_migration: $2432 \times 4 = 9728$

computation: 171892 (including write outActi from SRAM to SRAM)

"VGG-CONV_8": 413733

weight_clock: 19114

inActi_clock: 49389

weight_migration: $4736 \times 4 = 18944$

computation: 326286 (including write outActi from SRAM to SRAM)

"VGG-FC_19" : 132044

Result for VGG-19

MAC array Row = 2, Column = 16

"VGG-CONV_1" : 49153

weight_clock: 495

inActi_clock: 8374

weight_migration: 0

computation: 40284 (including write outActi from SRAM to SRAM)

"VGG-CONV_6" : 327914

weight_clock: 9611

inActi_clock: 23342

weight_migration: $2432 \times 4 = 9728$

computation: 285233 (including write outActi from SRAM to SRAM)

"VGG-CONV_8": 639999

weight_clock: 19115

inActi_clock: 49386

weight_migration: $4736 \times 4 = 18944$

computation: 552554 (including write outActi from SRAM to SRAM)

"VGG-FC_19" : 131864

Result for VGG-19

MAC array Row = 4, Column = 8

"VGG-CONV_1" : 52504

weight_clock: 493

inActi_clock: 8373

weight_migration: 0

computation: 43638 (including write outActi from SRAM to SRAM)

"VGG-CONV_6" : 376141

weight_clock: 9611

inActi_clock: 23342

weight_migration: $2433 \times 4 = 9732$

computation: 333456 (including write outActi from SRAM to SRAM)

"VGG-CONV_8": 736468

weight_clock: 19115

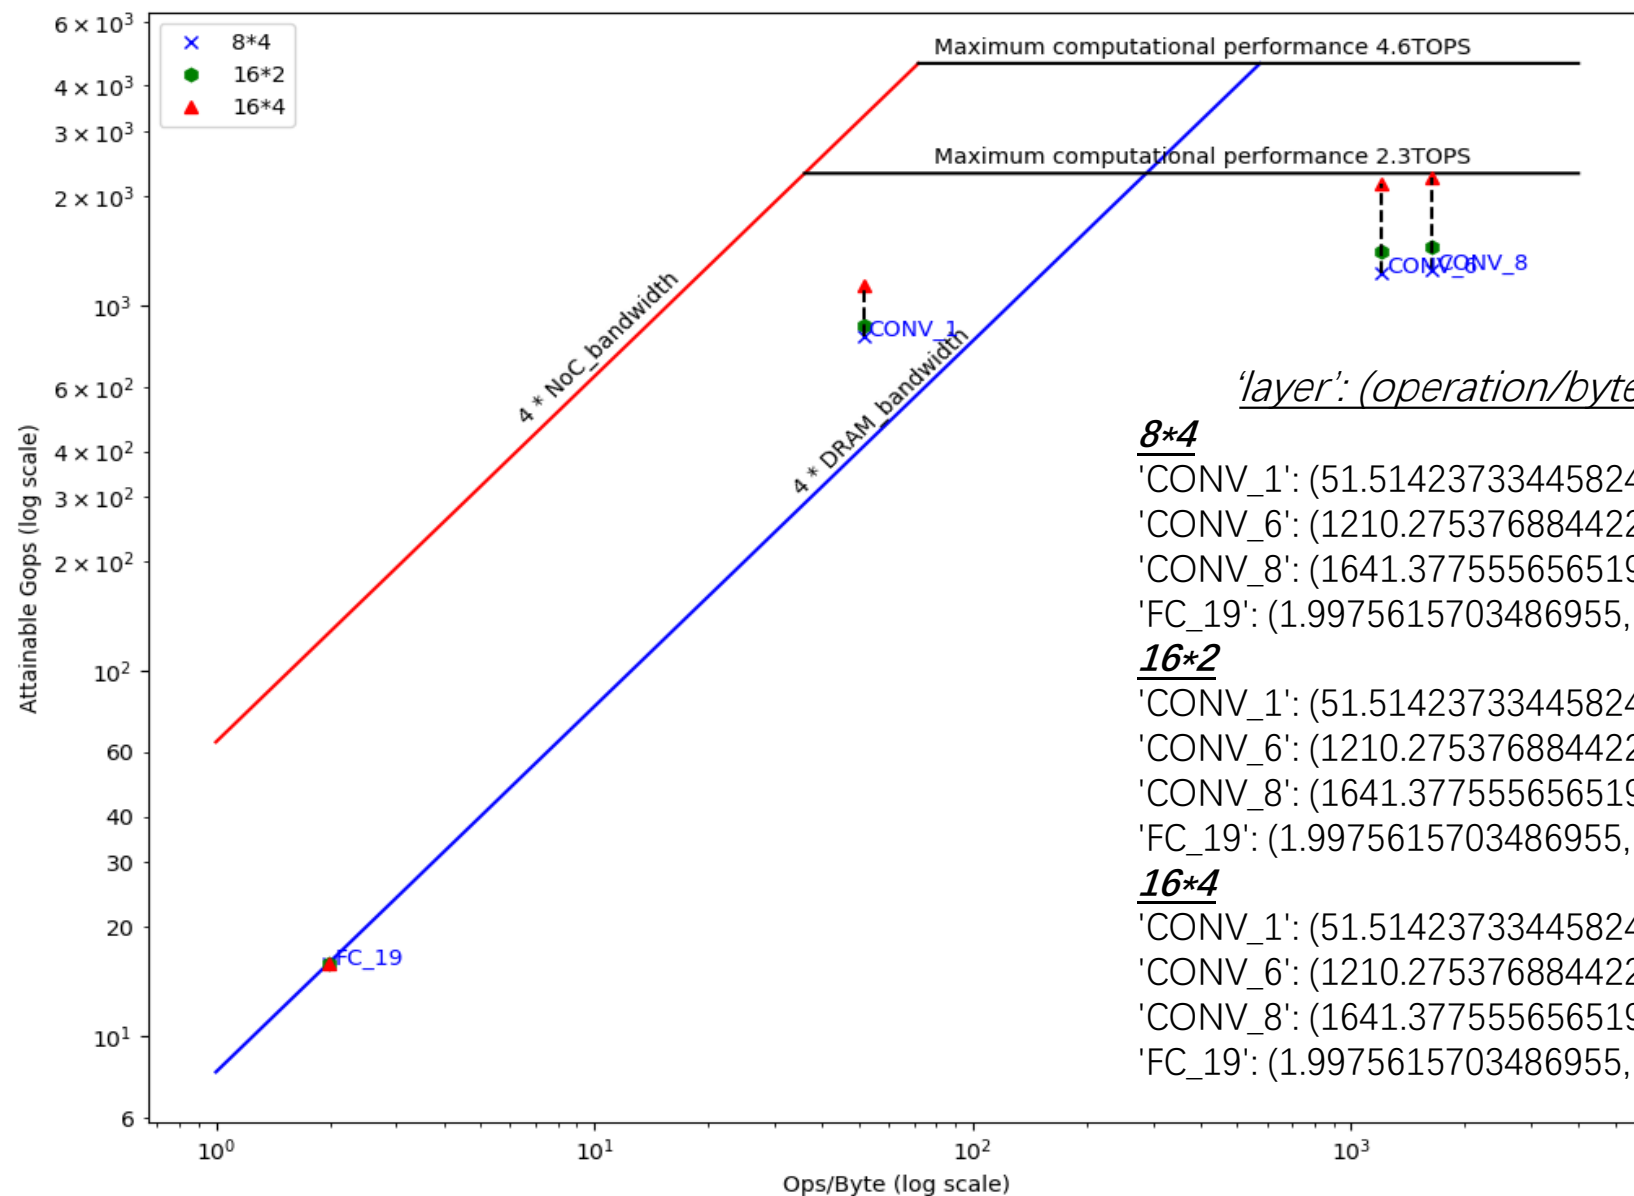
inActi_clock: 49388

weight_migration: $4736 \times 4 = 18944$

computation: 649021 (including write outActi from SRAM to SRAM)

"VGG-FC_19" : 132816

Result for VGG-19



'layer': (operation/bytes, attainable Gops)

8*4

'CONV_1': (51.51423733445824, 825.6906902331251)
 'CONV_6': (1210.275376884422, 1229.3847679460628)
 'CONV_8': (1641.3775556565197, 1255.783050994748)
 'FC_19': (1.9975615703486955, 15.789904830743284)

16*2

'CONV_1': (51.51423733445824, 881.9820560291334)
 'CONV_6': (1210.275376884422, 1410.192965228688)
 'CONV_8': (1641.3775556565197, 1445.0710579235279)
 'FC_19': (1.9975615703486955, 15.903900988897652)

16*4

'CONV_1': (51.51423733445824, 1141.3844452635458)
 'CONV_6': (1210.275376884422, 2155.0600768961904)
 'CONV_8': (1641.3775556565197, 2235.3644306835567)
 'FC_19': (1.9975615703486955, 15.882221077822544)

Result for ResNet-50

MAC array Row = 4, Column = 16

"RES-CONV_1": 99623
weight_clock: 971
inActi_clock: 11061
weight_migration: 0
computation: 87591 (including write outActi from SRAM to SRAM)

"CONV_SC_14": 115207
weight_clock: 6926
inActi_clock: 18284
weight_migration: 634 (half) * 2 = 1268
computation: 89363 (including write outActi from SRAM to SRAM)

"CONV_19": 44987
weight_clock: 7735
inActi_clock: 15453
weight_migration: 0
computation: 21799 (including write outActi from SRAM to SRAM)

"CONV_42": 47116
weight_clock: 25364
inActi_clock: 11049
weight_migration: 0
computation: 10703 (including write outActi from SRAM to SRAM)

Result for ResNet-50

MAC array Row = 2, , Column = 16

"RES-CONV_1":	168451 weight_clock: 968 inActi_clock: 11061 weight_migration: 0 computation: 156422 (including write outActi from SRAM to SRAM)
"CONV_SC_14":	180906 weight_clock: 6936 inActi_clock: 18290 weight_migration: 634 (half) * 2 = 1268 computation: 154412 (including write outActi from SRAM to SRAM)
"CONV_19":	59216 weight_clock: 7736 inActi_clock: 15455 weight_migration: 0 computation: 36025 (including write outActi from SRAM to SRAM)
"CONV_42":	55317 weight_clock: 25365 inActi_clock: 11051 weight_migration: 0 computation: 18901 (including write outActi from SRAM to SRAM)

Result for ResNet-50

MAC array Row = 4, Column = 8

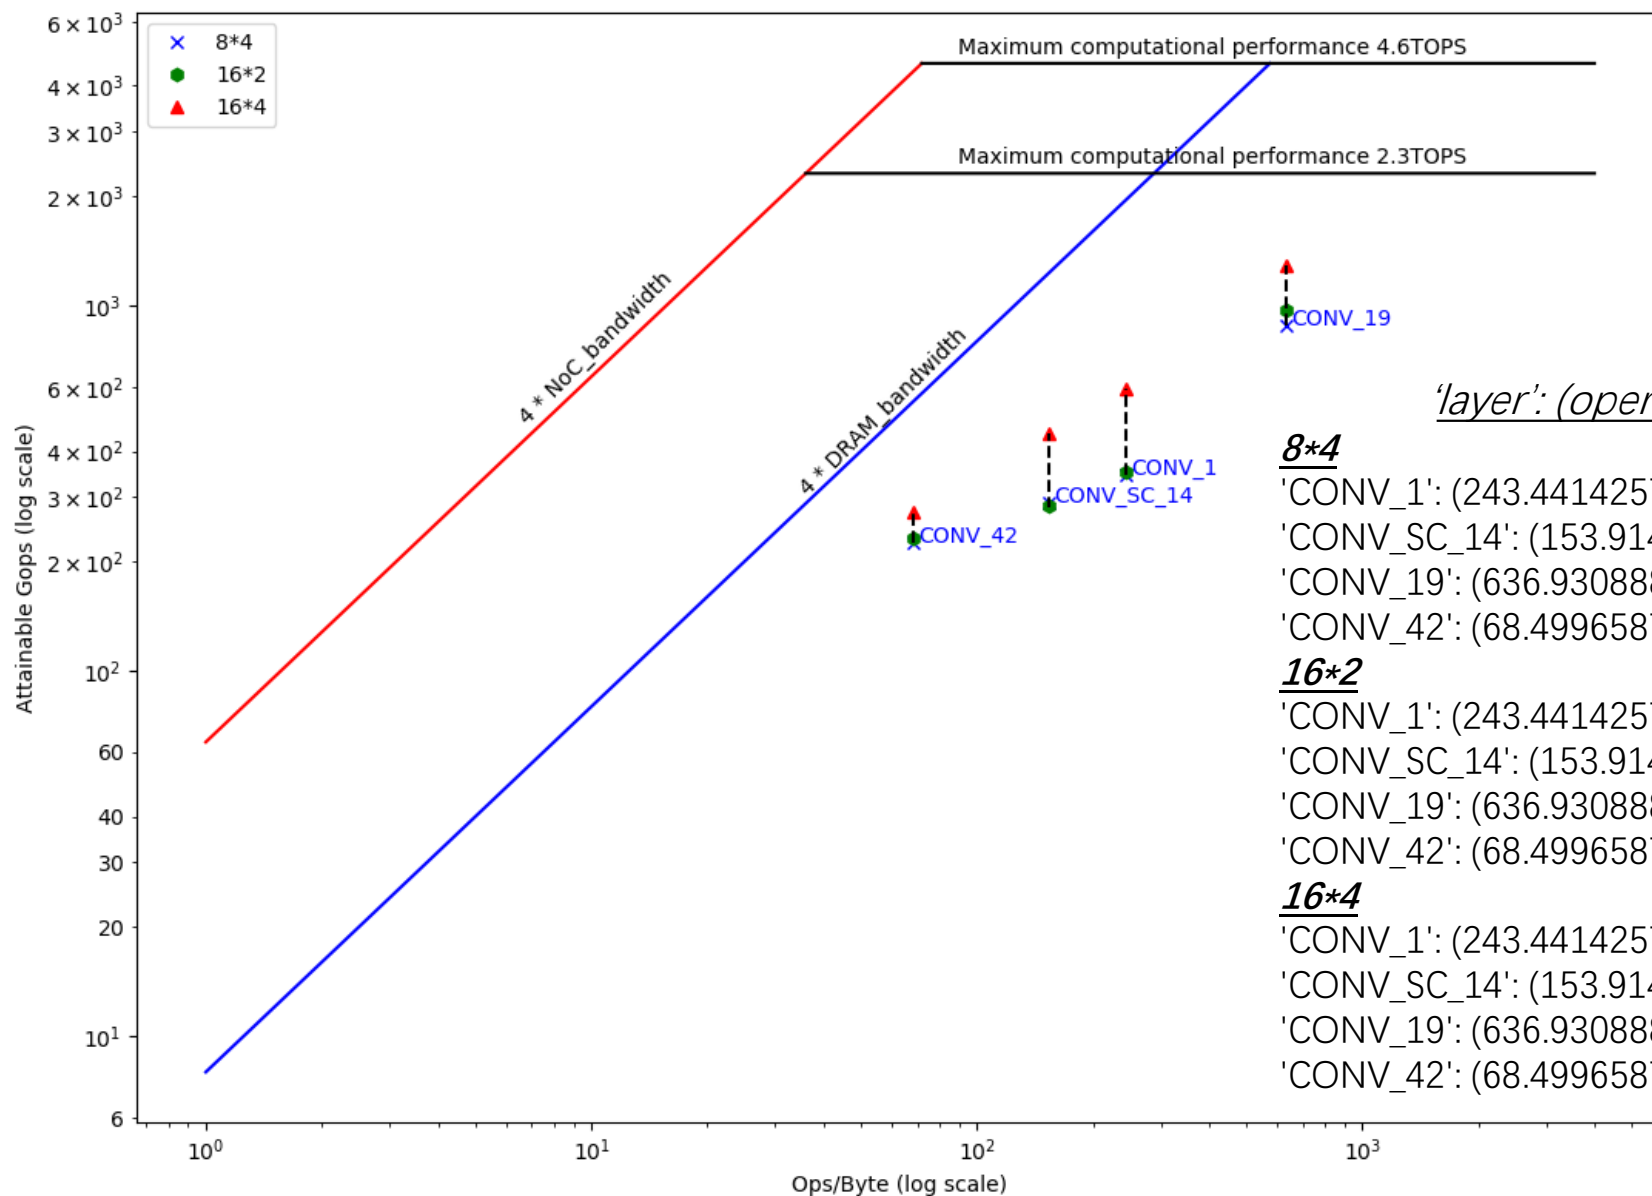
"RES-CONV_1": 170628
weight_clock: 969
inActi_clock: 11060
weight_migration: 0
computation: 158599 (including write outActi from SRAM to SRAM)

"CONV_SC_14": 176935
weight_clock: 6898
inActi_clock: 18295
weight_migration: 634 (half) * 2 = 1268
computation: 150474 (including write outActi from SRAM to SRAM)

"CONV_19": 65233
weight_clock: 7739
inActi_clock: 15454
weight_migration: 0
computation: 42040 (including write outActi from SRAM to SRAM)

"CONV_42": 57366
weight_clock: 25366
inActi_clock: 11050
weight_migration: 0
computation: 20950 (including write outActi from SRAM to SRAM)

Result for ResNet-50



Conclusion

- Decreasing computing resource to be half, the attainable computing power is not half. → Row=2 has higher computing resource utilization. It could better to balance the performance between different components. Or in other word, difference components match each other better.
- Decreasing computing resource to be half, the NoC's memory bandwidth problem is not so serious as Row=4. Current simulation is under $f_{NoC} = 2 * f_{pe}$. If $f_{NoC} < 2 * f_{pe}$, I think computing resource utilization of Row=2 is higher.
- Decreasing computing resource to be half, FC could get a little bit higher computing power, because the data amount increased by the alignment is decreased, which saving the times to output the result from MAC array to neighbor PE. However, It doesn't affect CONV.
- **16*2 is better than 8*4.**